

Paulo Eduardo dos Santos

Senior Big Data Engineer

Sao Paulo, Brazil — 100% Remote | paulo_eduardosp@yahoo.com.br | +55 11 91526-7337

LinkedIn: linkedin.com/in/paulo--eduardo | GitHub: github.com/Pauloeduspbr

PROFESSIONAL SUMMARY

Big Data Senior engineer with 9+ years as a developer and 8+ years in Big Data, writing highly efficient Apache Spark jobs that process terabytes of data and optimizing Spark workloads on large Hadoop clusters (Cloudera/Hortonworks) for banks and telecoms handling millions of daily transactions. Hands-on with batch and streaming (real-time) processing paradigms and large data loads with embedded data quality, operating resilient Big Data environments at 99.9% uptime SLA (Banco do Brasil) alongside production AWS platforms. Clean, efficient Python code; experienced mentor and technical lead in multicultural, distributed teams.

SKILLS

- **Big Data:** Apache Spark (TB-scale job optimization), PySpark, SparkSQL, Hadoop (Cloudera/Hortonworks), Hive, HDFS, Kafka, Impala
- **Languages:** Python (advanced), SQL (advanced), Shell Script, Scala
- **AWS & DevOps:** EMR / EMR Serverless, S3, Glue, Athena, Redshift, Lambda, Kinesis, CloudWatch; Docker, Kubernetes, Terraform, Ansible, CI/CD — resiliency, security, high availability (HA)
- **Data Quality & Loads:** data quality frameworks, large-scale batch/incremental loads, CDC, streaming ingestion (Kafka, Kinesis), observability (metrics, alerts, log handlers)
- **Analytics & Machine Learning Support:** pipelines feeding machine learning (ML) and analytics workloads — batch and streaming pipeline design, A/B test analysis with quality-assured datasets
- **Warehouses & Databases:** Snowflake, BigQuery, Redshift, PostgreSQL, Oracle, Cassandra, MongoDB

WORK EXPERIENCE

Data Engineer — BRQ (Clients: RD Saude, Itau)

Jul 2025 – Jun 2026 | Remote, Brazil

- Developed PySpark/SparkSQL pipelines ingesting and consolidating large-scale clinical data from 3 healthcare platforms (Amplified, HAOC, Vitat) on Apache Iceberg/S3 with Medallion architecture, orchestrated by Airflow over Spark on Kubernetes.
- Implemented end-to-end observability (custom log handlers, Spark metrics, exception tracking, AWS SNS + Step Functions alerts) ensuring pipeline resiliency and availability across dev/qa/prod Docker environments.
- Built data quality frameworks (cross-origin validations, light checks) for critical healthcare data loads.
- For Itau: built ETL pipelines with EMR Serverless and Glue Studio, optimized SQL for performance, and supported A/B test analysis with quality-assured datasets for hypothesis validation.

Senior Big Data Specialist / Data Engineer — Stefanini (Clients: Banco do Brasil, Cielo, AmericaNet, Itau, JBS)

Jun 2023 – Jun 2025 | Sao Paulo, Brazil

- Sustained and optimized production Cloudera/Hortonworks clusters processing millions of daily financial transactions at 99.9% uptime SLA (Banco do Brasil).
- Tuned Hive queries via partitioning and bucketing, and built a reusable Python automation framework that cut operational maintenance time by 40%.
- Developed PySpark scripts processing large transactional volumes for payment systems (Cielo), with S3 partitioning and external Athena tables; converted complex SQL logic to distributed Spark processing in Databricks.
- Processed terabyte-scale datasets with SparkSQL in Databricks and built ELT pipelines with dbt in Snowflake, reducing manual intervention by 70% through automation (AmericaNet).
- Refactored legacy stored procedures into PySpark jobs (Itau), achieving 60% storage reduction with Parquet.
- Led a team of engineers modernizing JBS data pipelines: high-availability NiFi cluster, Active Directory integration, load balancing, and S3-to-Redshift data movement.

Senior Data Engineer — Natura &Co

Feb 2022 – Jun 2023 | Sao Paulo, Brazil

- Built EMR Serverless ETL pipelines and complex Airflow DAGs with dynamic dependencies for a global e-commerce operation.

- Led the Parquet-to-Delta Lake migration with ACID transactions and schema evolution; incremental CDC loads integrated with BI via Redshift Spectrum.

- Modeled dimensional data for customer journey and cohort analysis across web analytics, transactional, and CRM sources.

Big Data Architect / Engineer — Hospital Sirio-Libanes

Mar 2019 – Dec 2020 | Sao Paulo, Brazil

- Built automated Airflow pipelines (Oracle extraction, S3 transfer, format conversion) on AWS (S3, Athena, EMR, Kinesis, Lambda).

- Containerized workloads with Docker, configured Kubernetes for scalability, and implemented SSO integrating AWS Cognito with Azure AD — security- and availability-oriented design.

- As Dell EMC consultant (Fleury): Hortonworks environment support, Kerberos/Ranger security hardening, Oozie job monitoring scripts, Grafana dashboards, and operational training for the client team.

Big Data Technical Lead — Claro Brasil

Sep 2017 – Mar 2019 | Sao Paulo, Brazil

- Implemented and operated 3 Cloudera clusters (Data Lake, M2M, BI); optimized heavy Spark/Hive workloads, cutting processing time by 50%.

- Secured clusters with Kerberos (AD as KDC) and cross-realm trust; integrated Hive/Impala with BI tools (SAS, QlikView, Alteryx).

- Mentored and led 2 engineers; built the data ingestion framework and monitoring/sustainment scripts.

Big Data Analyst — Vivo

Jun 2017 – Sep 2017 | Sao Paulo, Brazil

- Hadoop 2.2-to-2.6 migration, PySpark ETL pipelines, HBase persistence, Zookeeper-managed cluster HA.

EDUCATION

- **Associate Degree, Big Data and Data Intelligence** — Anhembí Morumbi University (expected 2026)

- **Artificial Intelligence** (in progress) — FMU

- **BSc, Information Systems** — UNIP (2006)

CERTIFICATIONS

- Red Hat Certified System Administrator (RHCSA)

- Big Data Engineering — Garre Training (160h)

LANGUAGES

- Portuguese (native) | English: Upper-Intermediate (B2) — fluent technical reading and writing